


# Exploring LLM-based Agents

[Jerry Zikun Chen](#)



[ChainML.net](#) / [Theoriq.ai](#)

Dec. 9th, 2024

# Traditional LLM Interaction



# What is an Agent?

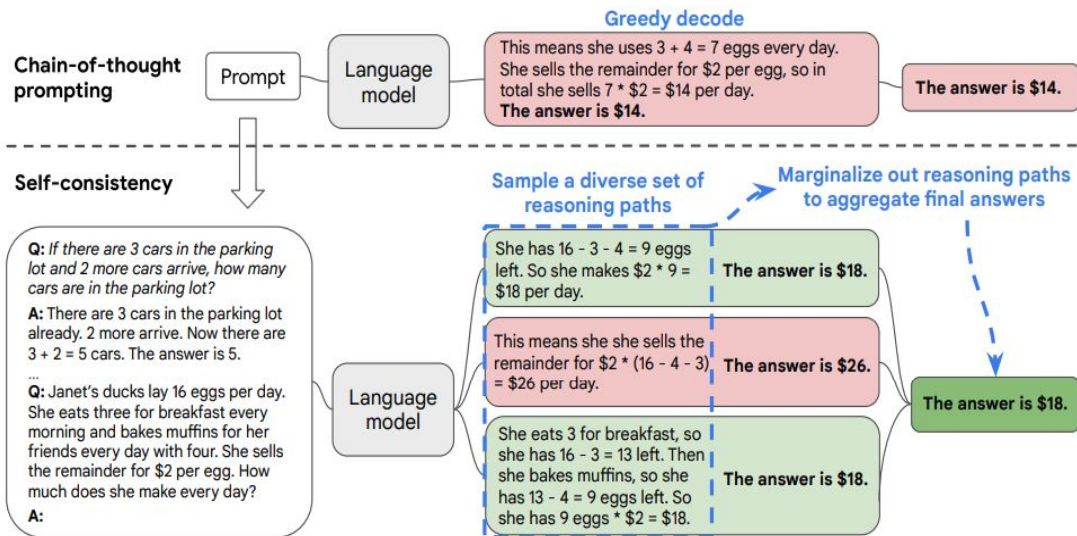


# Planning



- Subgoals: breaks down large complex tasks into smaller, more manageable subtasks (i.e. planning without feedback)
- Reflections: self-criticizes and reflects over past actions, learn from mistakes and refine for future iterations (i.e. planning with feedbacks)

# CoT and Self-Consistency



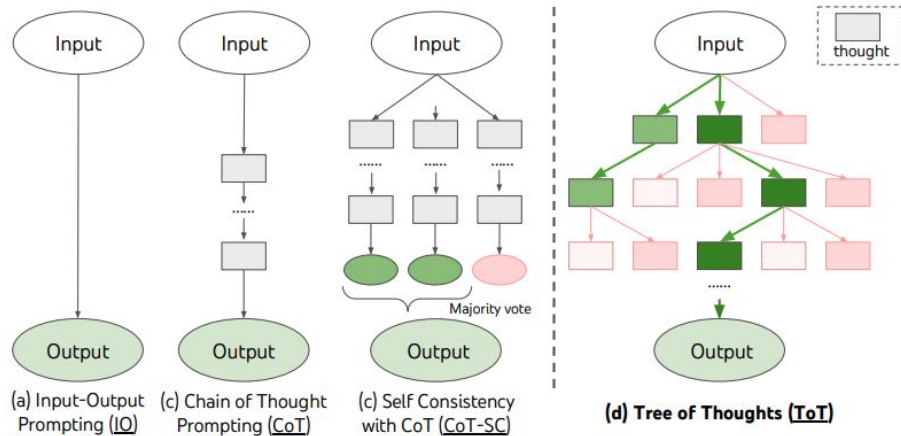
- **Chain of Thought**

- Instruct model to “Think step by step” to utilize more test-time compute to decompose complex tasks, making them more manageable
- Improves interpretability of model's internal thinking process

- **Self-Consistency**

- Replace the "greedy decoding" in CoT by sampling to generate a diverse set of reasoning paths
- Choose the most consistent answer (majority vote)

# ToT



- **Tree of Thought**

- Decomposes the problem into multiple thought steps
- Generates multiple thoughts per step, creating a tree structure
- Evaluate each state with a scalar value, classification / majority vote
- Use BFS or DFS to navigate the solution space
- Increase the capacity for complex reasoning by expanding and evaluating intermediate steps

# Planning



- Subgoals: breaks down large complex tasks into smaller, more manageable subtasks (i.e. planning without feedback)
- **Reflections: self-criticizes and reflects over past actions, learn from mistakes and refine for future iterations (i.e. planning with feedbacks)**

# ReAct: Synergizing Reasoning and Acting in Language Models

## (1d) ReAct (Reason + Act)

**Thought 1:** I need to search Apple Remote and find the program it was originally designed to interact with.

**Act 1:** `Search[Apple Remote]`

**Obs 1:** The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the `Front Row media center` program ...

**Thought 2:** Apple Remote was originally designed to control the Front Row media center program. I need to search `Front Row` next and find what other device can control it.

**Act 2:** `Search[Front Row]`

**Obs 2:** Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', ' `Front Row (software)` ', ...]

**Thought 3:** Front Row is not found. I need to search `Front Row (software)` .

**Act 3:** `Search[Front Row (software)]`

**Obs 3:** Front Row is a discontinued media center software ...

**Thought 4:** `Front Row (software)` is controlled by an Apple Remote or the keyboard function keys. So the answer is keyboard function keys.

**Act 4:** `Finish[keyboard function keys]`



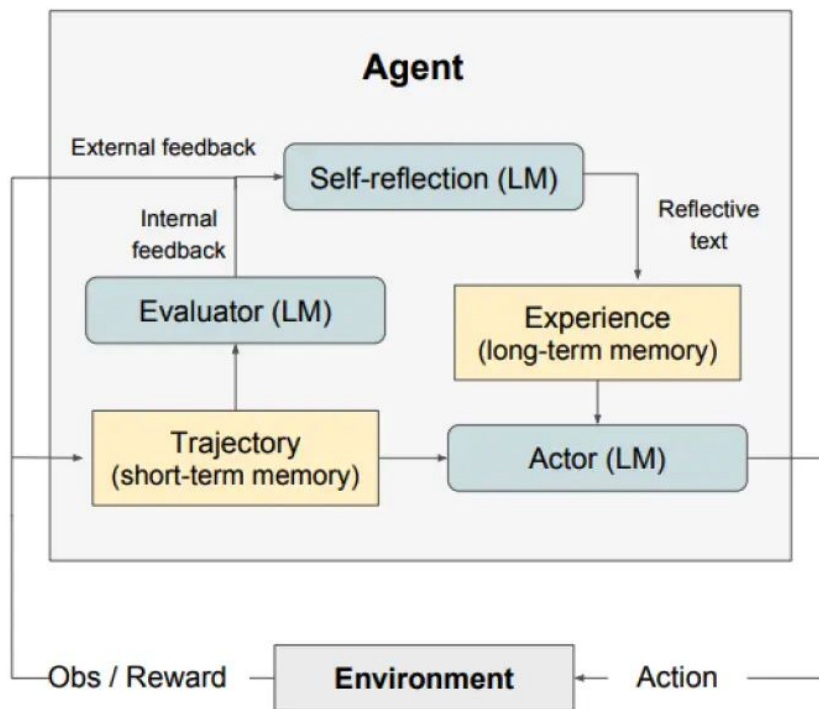
- Structured prompt with explicit step for LLM to think and act
  - Thought
  - Action
- Observations from environment serve as feedback to the agent
- Integrates reasoning and acting: Action space = task-specific discrete actions + language space
- Interleaving thoughts and actions
- Combines internal reasoning and external feedbacks



# Memory

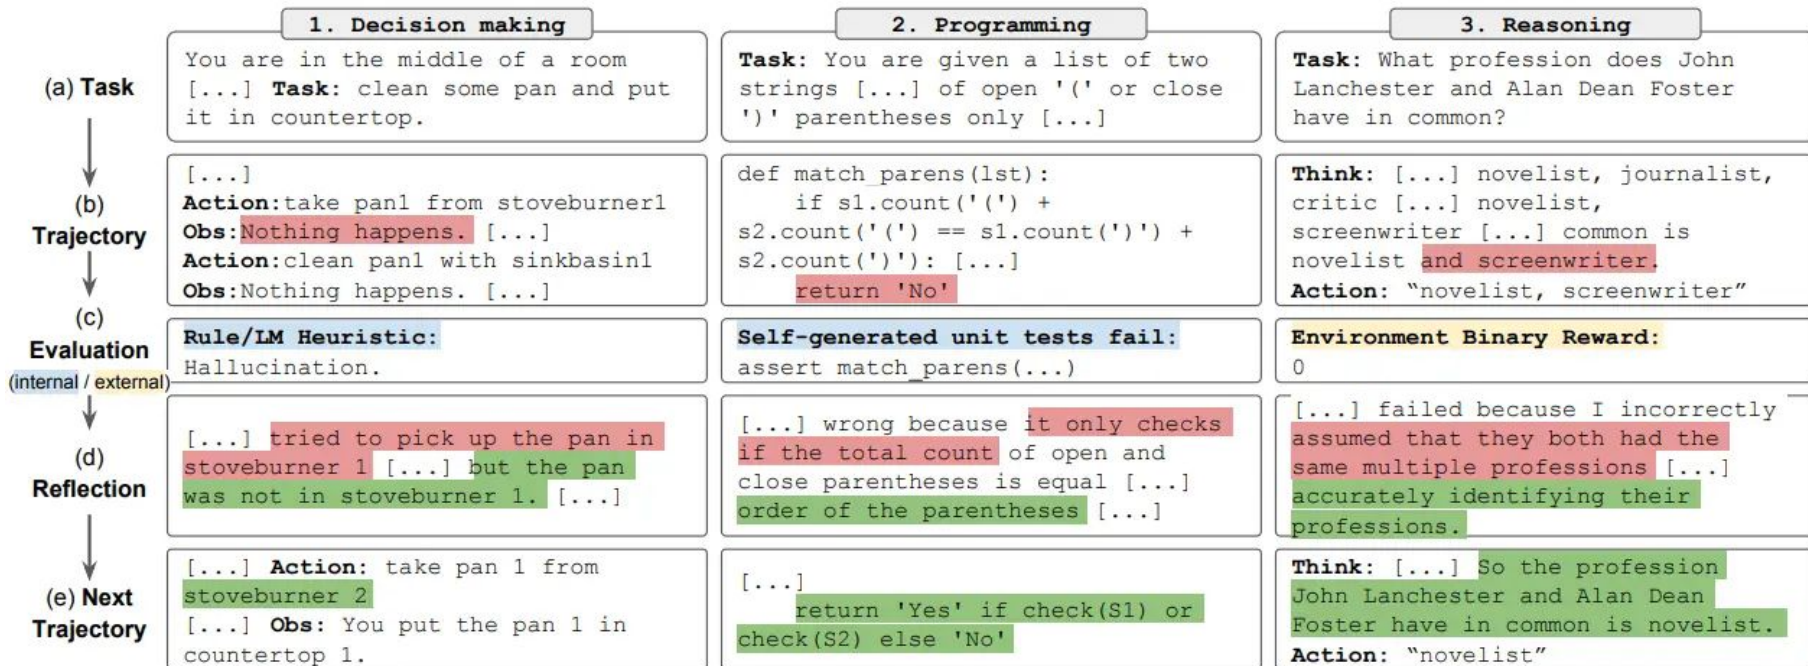
- Short-term / working memory: in-context learning / prompt engineering
  - finite due to context window constraints
- Long-term memory: agent can retain and recall information over extended periods, leveraging external vector DB and fast retrieval

# Reflexion






- Converts feedback (language or scalar) from the environment into linguistic feedback (self-reflection), provided as context for the agent in the next episode
- **Core components**
  - **Actor:** CoT, ReAct
  - **Evaluator:** trajectory -> reward score
  - **Self Reflection:** generate verbal reinforcement cues to improve the Actor
  - **Memory:**
    - short-term: current trajectory history
    - long-term: outputs from self-reflection (lessons learned over several trials)

# Reflexion



# Reflexion



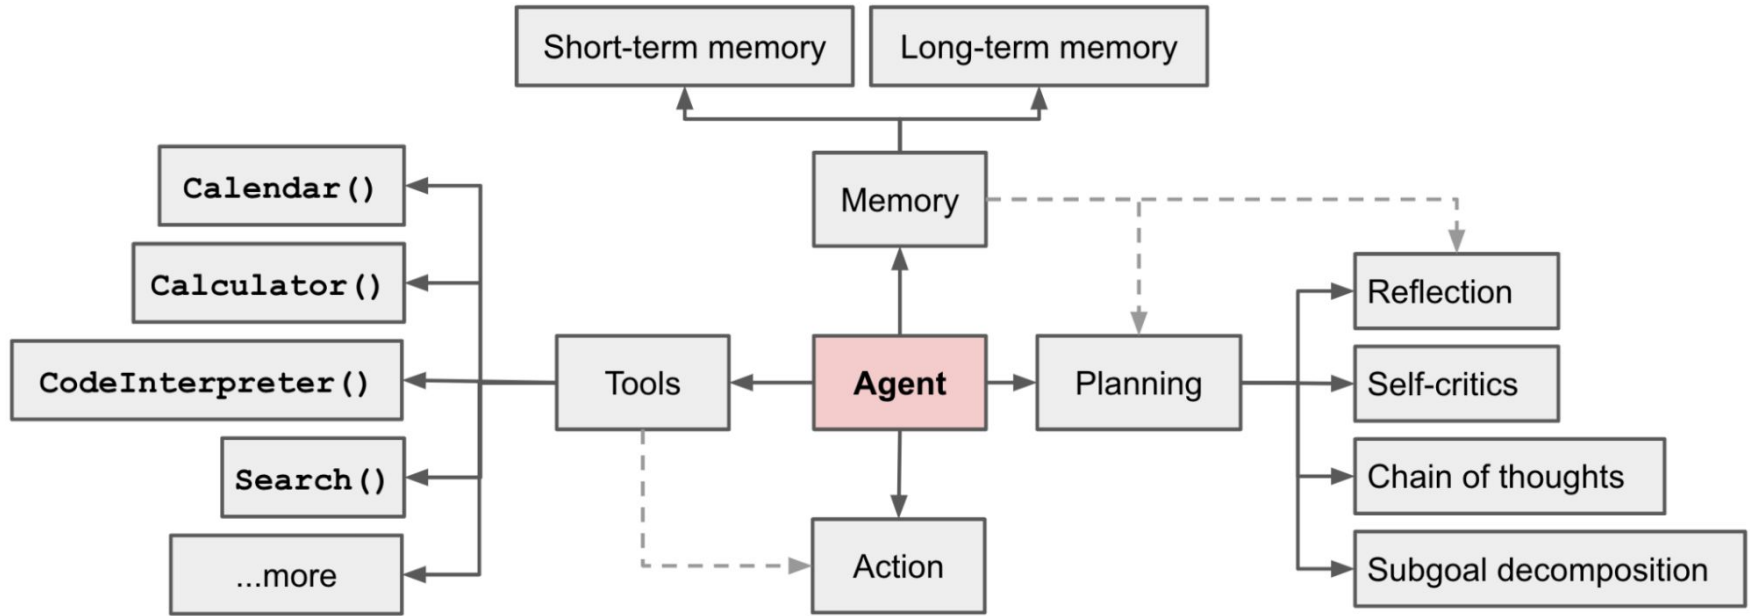
Advantages:

- The agent can learn from trial and error, reflecting on past mistakes for future decisions. e.g. decision-making, reasoning, programming etc.
- A lightweight alternative to traditional RL, which requires training data and expensive model parameter fine-tuning
- Verbal feedback is more nuanced and specific than a scalar reward in traditional RL, helping agent to better understand its mistakes
- Memory containing the reflections is more explicit and interpretable for understanding where the learning comes from


# Tool Use

- a.k.a. function calling
- allows an agent to perform actions, interacting with the external environments such as using:
  - search API
  - code interpreter
  - math engine
  - knowledge bases
  - github access
  - ...

# Agent in a Nutshell



# Single Agent Challenges




- An agent can have **too many tools** to choose from and make poor decision about which tool to call next
- **Context grows too complex** for single agent to keep track of (e.g. too many reflections of experiences with different tasks)
- There can be a need for multiple specialization areas in the system (e.g. planner, researcher, code expert etc.) to **make optimization of individual capabilities easier**

# Why Multi-Agent Systems?


- Modularity
  - Separate agents are easier to develop, test, and maintain
- Specialization
  - Create experts focused on specific domains that helps the overall system performance
- Control
  - Explicitly control how agents communicate, compared to relying on an agent's decision on function and tool calling



# Multi-Agent Systems



# Multi-Agent Systems



# Multi-Agent Systems History

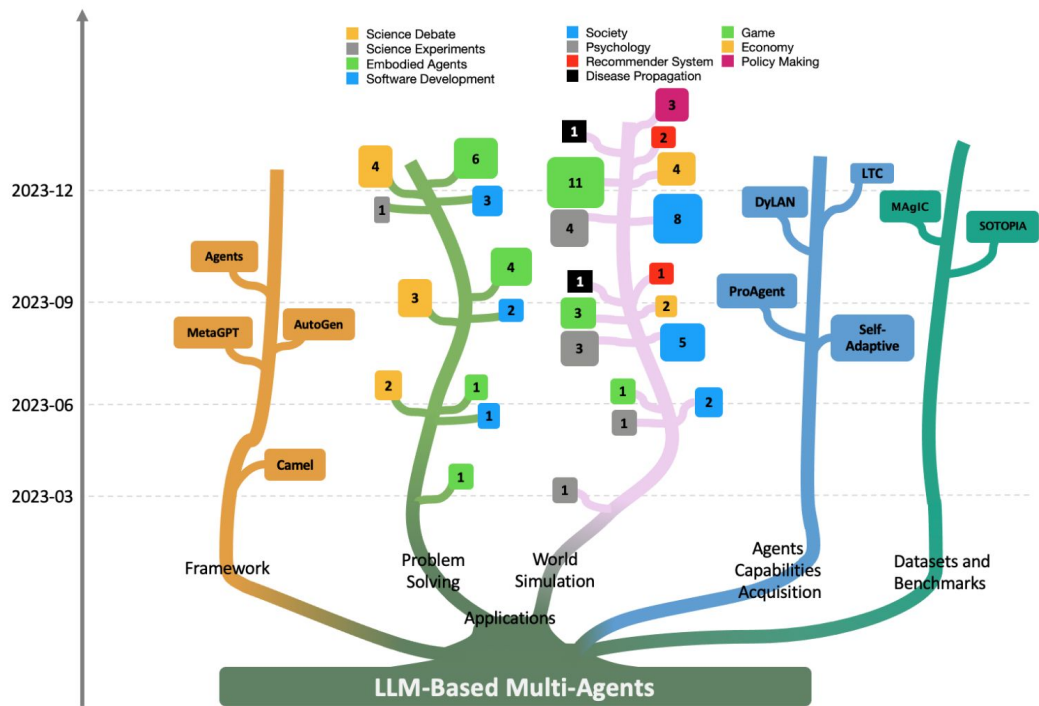


Figure 1: The rising trend in the research field of LLM-based Multi-Agents. For Problem Solving and World Simulation, we categorize current work into several categories and count the number of papers of different types at 3-month intervals. The number at each leaf node denotes the count of papers within that category.

# Aspects of Multi-Agent Systems

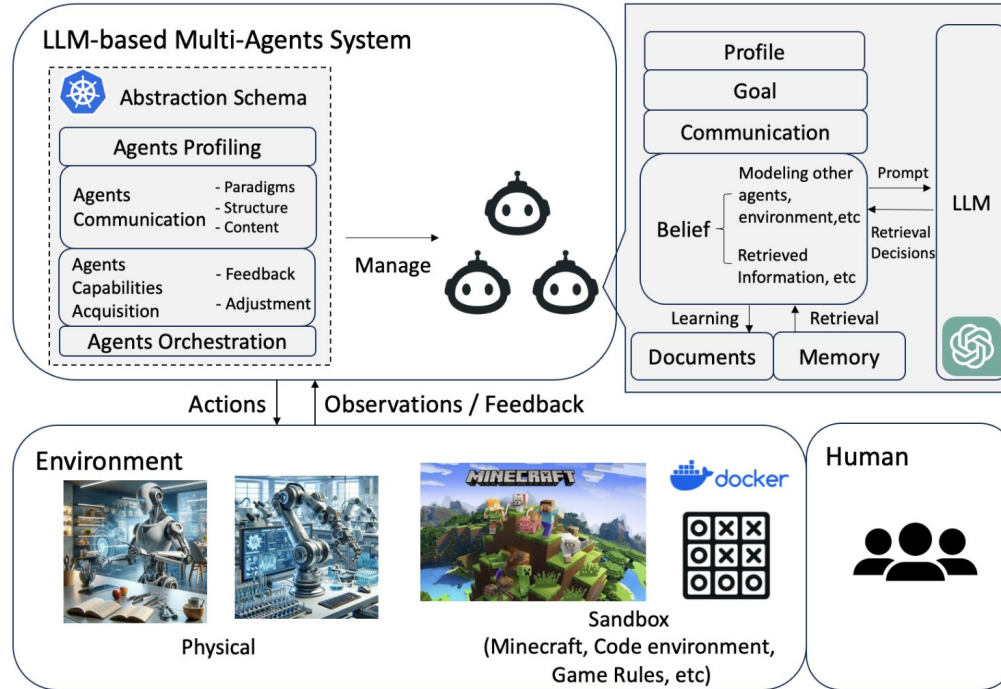


Figure 2: The Architecture of LLM-MA Systems.

# Multi-Agent Communication

- **Cooperative:** agents work together towards a shared objective, and exchange information to enhance a collective solution
- **Debate:** agents engage in argumentative interactions, presenting and defending their own viewpoints or solutions, and critiquing others
- **Competitive:** agents working towards their own goals that might be in conflict with goals of other agents


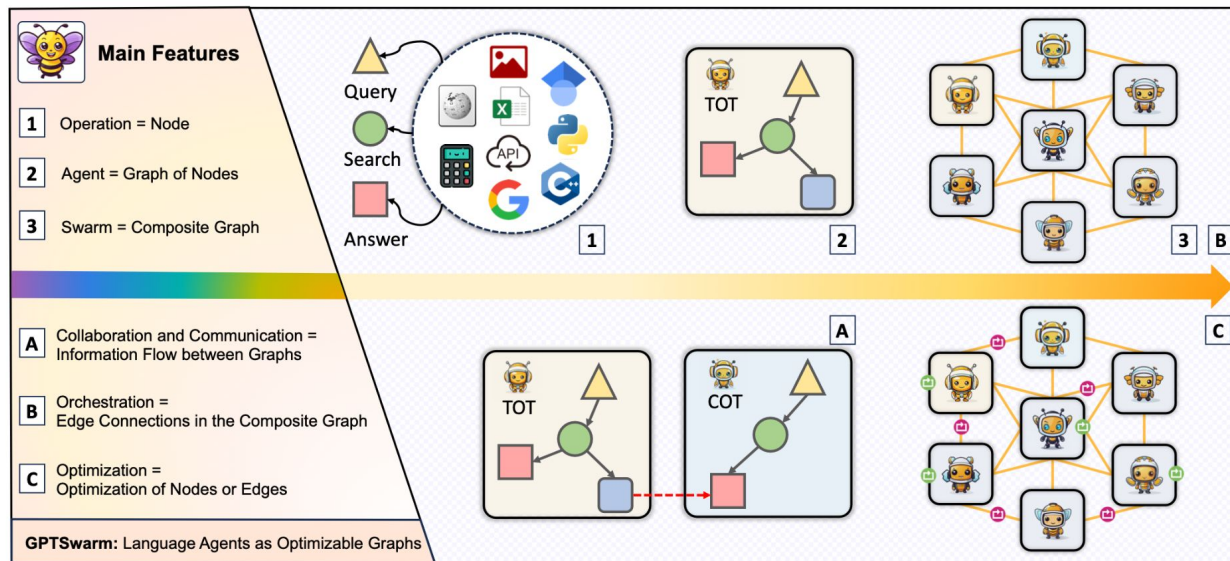


Figure 3: The Agent Communication Structure.

# Agents as Graphs

## GPTSwarm: Language Agents as Optimizable Graphs



**Figure 1. GPTSwarm is a framework that represents agents as graphs.** In this framework, each node represents an operation (e.g., LLM inference or tool use). An agent is a graph composed of these nodes. An edge between two agent graphs characterizes a communication channel; each agent collaborates with others through different channels. When connected, multiple agents form a composite graph with a certain orchestration topology. This graph representation lends itself to optimization of nodes and edges via prompting and evolutionary or reinforcement learning techniques.

# Agent Frameworks

- CrewAI
- LangGraph from LangChain
- Autogen from Microsoft
- OpenAI Swarm

# CrewAI

## **Pros:**

- Straightforward and beginner-friendly setup
- Built on top of LangChain, an established framework with lots of of community support, documentation, and useful tools
- Clear documentation with working examples and YouTube tutorials

## **Cons:**

- Customizing things can get more complex
- Reliance on LangChain adds another dependency to your project



# LangGraph

## Pros:

- Low-level framework that gives you fine-grained control over each element, making it powerful for complex workflows
- Seamless integration with the LangChain ecosystem (in particular, the very large number of available built-in tools)
- Checkpoints can capture the past and present states of the agent for monitoring and error recovery purposes
- LangGraph Studio, a new specialized IDE, lets you visualize, interact with, and debug your agent workflows

## Cons:

- Documentation can be hit or miss, sometimes outdated or missing key details
- Not the easiest framework to work with initially

# Microsoft Autogen

## Pros:

- Agents can generate, fix, and run code in Docker containers
- Autogen Studio, a no-code tool, making it easier to manage and debug
- Great option if you like a visual approach
- Offers a lot of flexibility and plenty of conversation patterns to work with
- Clear documentation with useful examples

## Cons:

- Works best with the latest models; older ones may struggle with tasks
- May have a bit of a learning curve

# OpenAI Swarm

## **Pros:**

- Clean and easy-to-work-with structure, making it perfect for experimentation
- Simplest, cleanest, and most lightweight of the options

## **Cons:**

- Still experimental and not meant for production use
- Does not work with the assistance API
- Limited features and capabilities compared to other frameworks
- Not a lot of community backing or support

# References

## Papers

Tree of Thoughts: <https://arxiv.org/abs/2305.10601>

ReAct: <https://arxiv.org/abs/2210.03629>

Reflexion: <https://arxiv.org/abs/2303.11366>

LLM-Based Agents Survey:  
<https://arxiv.org/pdf/2402.01680>

GPTSwarm: <https://arxiv.org/pdf/2402.16823>

## YouTube

[https://www.youtube.com/watch?v=pEMhPBQMNjg&ab\\_channel=SamWitteveen](https://www.youtube.com/watch?v=pEMhPBQMNjg&ab_channel=SamWitteveen)

[https://www.youtube.com/watch?v=hvAPnpSfSGo&t=1s&ab\\_channel=LangChain](https://www.youtube.com/watch?v=hvAPnpSfSGo&t=1s&ab_channel=LangChain)

[https://www.youtube.com/watch?v=4nZI32FwU-o&ab\\_channel=LangChain](https://www.youtube.com/watch?v=4nZI32FwU-o&ab_channel=LangChain)

## Others

[https://www.youtube.com/watch?v=q1XFm21l-VQ&t=369s&ab\\_channel=SnowflakeDevelopers](https://www.youtube.com/watch?v=q1XFm21l-VQ&t=369s&ab_channel=SnowflakeDevelopers)

<https://lilianweng.github.io/posts/2023-06-23-agent/>

<https://www.promptingguide.ai/research/llm-agents>

<https://developer.nvidia.com/blog/introduction-to-llm-agents/>

<https://blog.dataiku.com/open-source-frameworks-for-llm-powered-agents>

# LangGraph Demo